

11. Worksheet: Phylogenetic Diversity - Traits

Lana Bolin; Z620: Quantitative Biodiversity, Indiana University

19 February, 2019

OVERVIEW

Up to this point, we have been focusing on patterns taxonomic diversity in Quantitative Biodiversity. Although taxonomic diversity is an important dimension of biodiversity, it is often necessary to consider the evolutionary history or relatedness of species. The goal of this exercise is to introduce basic concepts of phylogenetic diversity.

After completing this exercise you will be able to:

1. create phylogenetic trees to view evolutionary relationships from sequence data
2. map functional traits onto phylogenetic trees to visualize the distribution of traits with respect to evolutionary history
3. test for phylogenetic signal within trait distributions and trait-based patterns of biodiversity

Directions:

1. In the Markdown version of this document in your cloned repo, change “Student Name” on line 3 (above) with your name.
2. Complete as much of the worksheet as possible during class.
3. Use the handout as a guide; it contains a more complete description of data sets along with examples of proper scripting needed to carry out the exercises.
4. Answer questions in the worksheet. Space for your answers is provided in this document and is indicated by the “>” character. If you need a second paragraph be sure to start the first line with “>”. You should notice that the answer is highlighted in green by RStudio (color may vary if you changed the editor theme).
5. Before you leave the classroom today, it is *imperative* that you **push** this file to your GitHub repo, at whatever stage you are. This will enable you to pull your work onto your own computer.
6. When you have completed the worksheet, **Knit** the text and code into a single PDF file by pressing the **Knit** button in the RStudio scripting panel. This will save the PDF output in your ‘8.BetaDiversity’ folder.
7. After Knitting, please submit the worksheet by making a **push** to your GitHub repo and then create a **pull request** via GitHub. Your pull request should include this file (**11.PhyloTraits_Worksheet.Rmd**) with all code blocks filled out and questions answered) and the PDF output of Knitr (**11.PhyloTraits_Worksheet.pdf**).

The completed exercise is due on **Wednesday, February 20th, 2019 before 12:00 PM (noon)**.

1) SETUP

Typically, the first thing you will do in either an R script or an RMarkdown file is setup your environment. This includes things such as setting the working directory and loading any packages that you will need.

In the R code chunk below, provide the code to:

1. clear your R environment,
2. print your current working directory,
3. set your working directory to your “/11.PhyloTraits” folder, and
4. load all of the required R packages (be sure to install if needed).

```

rm(list = ls())
getwd()

## [1] "/Users/lana/GitHub/QB2019_Bolin/2.Worksheets/11.PhyloTraits"
setwd("/Users/lana/GitHub/QB2019_Bolin/2.Worksheets/11.PhyloTraits")

package.list <- c('ape', 'seqinr', 'phylobase', 'adephylo', 'geiger', 'picante', 'stats', 'RColorBrewer')
for (package in package.list) {
  if (!require(package, character.only=TRUE, quietly=TRUE)) {
    install.packages(package)
    library(package, character.only=TRUE)
  }
}

##
## Attaching package: 'seqinr'
## The following objects are masked from 'package:ape':
##
##   as.alignment, consensus
##
## Attaching package: 'phylobase'
## The following object is masked from 'package:ape':
##
##   edges
##
## Attaching package: 'permute'
## The following object is masked from 'package:seqinr':
##
##   getType
## This is vegan 2.5-3
##
## Attaching package: 'nlme'
## The following object is masked from 'package:seqinr':
##
##   gls
##
## Attaching package: 'dplyr'
## The following object is masked from 'package:MASS':
##
##   select
## The following object is masked from 'package:nlme':
##
##   collapse
## The following object is masked from 'package:seqinr':
##
##   count
## The following objects are masked from 'package:stats':

```

```
##
##      filter, lag
## The following objects are masked from 'package:base':
##
##      intersect, setdiff, setequal, union
##
## Attaching package: 'phangorn'
## The following objects are masked from 'package:vegan':
##
##      diversity, treedist
```

2) DESCRIPTION OF DATA

The maintenance of biodiversity is thought to be influenced by **trade-offs** among species in certain functional traits. One such trade-off involves the ability of a highly specialized species to perform exceptionally well on a particular resource compared to the performance of a generalist. In this exercise, we will take a phylogenetic approach to mapping phosphorus resource use onto a phylogenetic tree while testing for specialist-generalist trade-offs.


3) SEQUENCE ALIGNMENT

Question 1: Using your favorite text editor, compare the `p.isolates.fasta` file and the `p.isolates.afa` file. Describe the differences that you observe between the two files.

Answer 1: The `.afa` file has deletions marked as dashes, while the `.fasta` file doesn't.

In the R code chunk below, do the following: 1. read your alignment file, 2. convert the alignment to a DNABin object, 3. select a region of the gene to visualize (try various regions), and 4. plot the alignment using a grid to visualize rows of sequences.

```
read.aln <- read.alignment(file = "./data/p.isolates.afa", format = "fasta")
p.DNABin <- as.DNABin(read.aln)
window <- p.DNABin[, 100:500]
image.DNABin(window, cex.lab = 0.5)
```



Question 2: Make some observations about the `muscle` alignment of the 16S rRNA gene sequences for our bacterial isolates and the outgroup, *Methanosarcina*, a member of the domain Archaea. Move along the alignment by changing the values in the `window` object.

- Approximately how long are our sequence reads?
- What regions do you think would be appropriate for phylogenetic inference and why?

Answer 2a: ~400 bp **Answer 2b:** Regions that show a lot of variation among isolates are appropriate because that allows those regions to be grouped in a meaningful way (in other words, if we chose regions without much variation our tree would have lots of unhelpful polytomies).

4) MAKING A PHYLOGENETIC TREE

Once you have aligned your sequences, the next step is to construct a phylogenetic tree. Not only is a phylogenetic tree effective for visualizing the evolutionary relationship among taxa, but as you will see later, the information that goes into a phylogenetic tree is needed for downstream analysis.

A. Neighbor Joining Trees

In the R code chunk below, do the following:

- calculate the distance matrix using `model = "raw"`,
- create a Neighbor Joining tree based on these distances,
- define “*Methanosarcina*” as the outgroup and root the tree, and
- plot the rooted tree.

```

# 1
seq.dist.raw <- dist.dna(p.DNAbin, model = "raw", pairwise.deletion = FALSE)

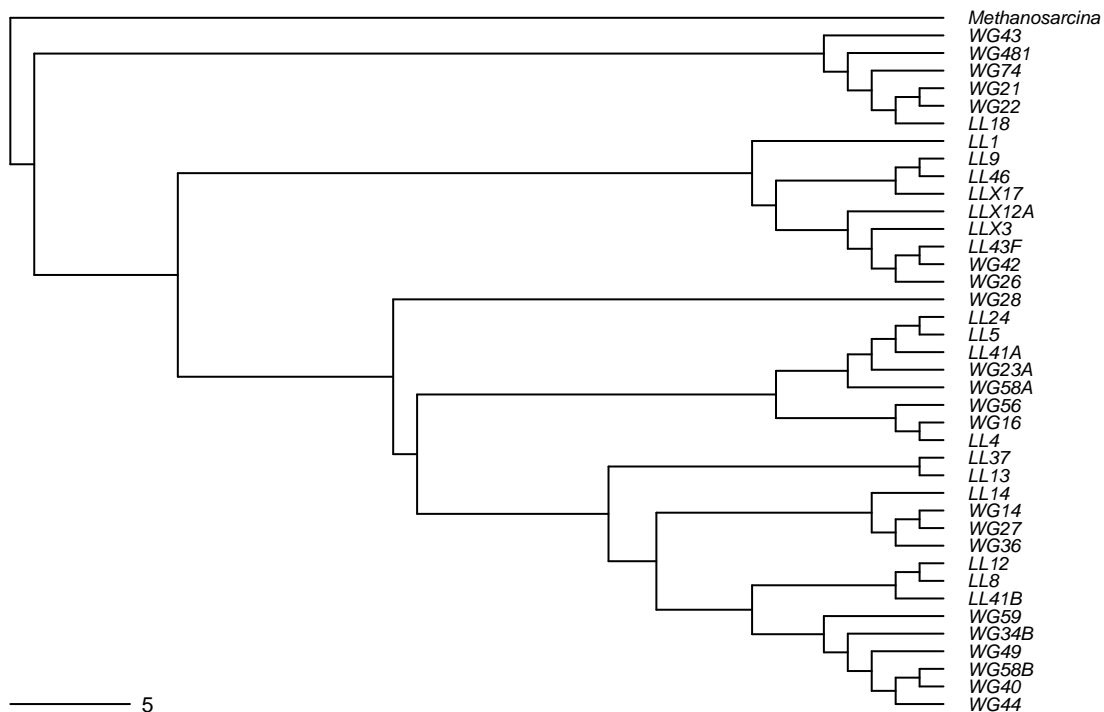
# 2
nj.tree <- bionj(seq.dist.raw)

# 3
outgroup <- match("Methanosarcina", nj.tree$tip.label)

# 4
nj.rooted <- root(nj.tree, outgroup, resolve.root = TRUE)
par(mar = c(1, 1, 2, 1) + 0.1)
plot.phylo(nj.rooted, main = "Neighbor Joining Tree", "phylogram", use.edge.length = FALSE, direction =
add.scale.bar(cex = 0.7)

```

Neighbor Joining Tree



Question 3: What are the advantages and disadvantages of making a neighbor joining tree?

Answer 3: It's fast so it's a great place to start looking at your data, but it's not as sophisticated as some other types of tree. For example, it doesn't take into account basepair transition probabilities.

B) SUBSTITUTION MODELS OF DNA EVOLUTION


In the R code chunk below, do the following:

1. make a second distance matrix based on the Felsenstein 84 substitution model,
2. create a saturation plot to compare the *raw* and *Felsenstein (F84)* substitution models,

3. make Neighbor Joining trees for both, and
4. create a cophylogenetic plot to compare the topologies of the trees.

```
# 1
seq.dist.F84 <- dist.dna(p.DNABin, model = "F84", pairwise.deletion = FALSE)

# 2
par(mar = c(5, 5, 2, 1) + 0.1)
plot(seq.dist.raw, seq.dist.F84,
     pch = 20, col = "red", las = 1, asp = 1, xlim = c(0, 0.7), ylim = c(0, 0.7), xlab = "Raw Distance",
     abline(b = 1, a = 0, lty = 2)
     text(0.65, 0.6, "1:1"))
```




```
# 3
raw.tree <- bionj(seq.dist.raw)
F84.tree <- bionj(seq.dist.F84)

raw.outgroup <- match("Methanosarcina", raw.tree$tip.label)
F84.outgroup <- match("Methanosarcina", F84.tree$tip.label)

raw.rooted <- root(raw.tree, raw.outgroup, resolve.root = TRUE)
F84.rooted <- root(F84.tree, F84.outgroup, resolve.root = TRUE)

# 4
layout(matrix(c(1, 2), 1, 2), width = c(1, 1))
par(mar = c(1, 1, 2, 0))
plot.phylo(raw.rooted, type = "phylogram", direction = "right", show.node.label = TRUE, use.edge.length = TRUE)
```

```
par(mar = c(1, 0, 2, 1))
plot.phylo(F84.rooted, type = "phylogram", direction = "left", show.node.label = TRUE, use.edge.length = TRUE)
```




In the R code chunk below, do the following:

1. pick another substitution model,
2. create a distance matrix and tree for this model,
3. make a saturation plot that compares that model to the *Felsenstein (F84)* model,
4. make a cophylogenetic plot that compares the topologies of both models, and
5. be sure to format, add appropriate labels, and customize each plot.

```
# 1: Tamura model

# 2
seq.dist.T92 <- dist.dna(p.DNABin, model = "T92", pairwise.deletion = FALSE)


# 3
par(mar = c(5, 5, 2, 1) + 0.1)
plot(seq.dist.T92, seq.dist.F84,
     pch = 20, col = "red", las = 1, asp = 1, xlim = c(0, 0.7), ylim = c(0, 0.7), xlab = "T92 Distance")
abline(b = 1, a = 0, lty = 2)
text(0.65, 0.6, "1:1")
```



```
# 4
T92.tree <- bionj(seq.dist.T92)
T92.outgroup <- match("Methanosarcina", T92.tree$tip.label)
T92.rooted <- root(T92.tree, T92.outgroup, resolve.root = TRUE)

layout(matrix(c(1, 2), 1, 2), width = c(1, 1))
par(mar = c(1, 1, 2, 0))
plot.phylo(T92.rooted, type = "phylogram", direction = "right", show.node.label = TRUE, use.edge.length = TRUE)

par(mar = c(1, 0, 2, 1))
plot.phylo(F84.rooted, type = "phylogram", direction = "left", show.node.label = TRUE, use.edge.length = TRUE)
```

Question 4:

- Describe the substitution model that you chose. What assumptions does it make and how does it compare to the F84 model?
- Using the saturation plot and cophylogenetic plots from above, describe how your choice of substitution model affects your phylogenetic reconstruction. If the plots are inconsistent with one another, explain why.
- How does your model compare to the *F84* model and what does this tell you about the substitution rates of nucleotide transitions?

Answer 4a: I chose the Tamura model (T92), which assumes equal frequencies of nucleotides but recognizes that transition mutations occur with higher probability than transversion mutations, and accounts for G + C content of the sequences. **Answer 4b:** T92 and F84 models create the same phylogeny. Their saturation plot has points along the 1:1 line, and the trees match perfectly. **Answer 4c:** Because these models create identical trees, substitution rates of nucleotide transitions are not different enough to change the tree structure.

C) ANALYZING A MAXIMUM LIKELIHOOD TREE

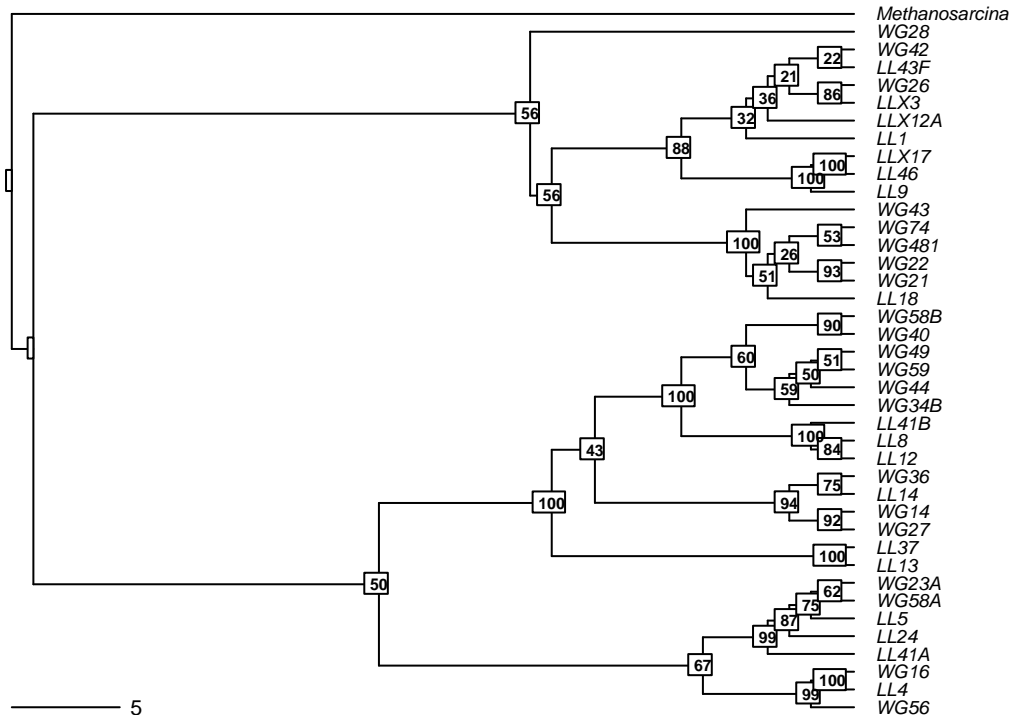
In the R code chunk below, do the following:

- Read in the maximum likelihood phylogenetic tree used in the handout.
- Plot bootstrap support values onto the tree

```
ml.bootstrp <- read.tree("./data/ml_tree/RAxML_bipartitions.T1")
par(mar = c(1, 1, 2, 1), + 0.1)
plot.phylo(ml.bootstrp, type = "phylogram", direction = "right", show.tip.label = TRUE, use.edge.length = TRUE)
```

```
add.scale.bar(cex = 0.7)
nodelabels(ml.bootstrap$node.label, font = 2, bg = "white", frame = "r", cex = 0.5)
```

Maximum Likelihood with Support Values



Question 5:

- How does the maximum likelihood tree compare the to the neighbor-joining tree in the handout? If the plots seem to be inconsistent with one another, explain what gives rise to the differences.
- Why do we bootstrap our tree?
- What do the bootstrap values tell you?
- Which branches have very low support?
- Should we trust these branches?

Answer 5a: It's quite different from the neighbor-joining tree, presumably because it takes into account nucleotide states? **Answer 5b:** Bootstrapping lets us run stats on our and generate confidence intervals. **Answer 5c:** They tell you in what proportion of your bootstrapped trees did a particular node/branch show up. A high value indicates the node/branch is well supported. **Answer 5d:** Many of the branches in the clade that includes WG42, LL43F, WG26, LLX3, LLX12A, and LL1 have very low support. **Answer 5e:** No, we should not.

5) INTEGRATING TRAITS AND PHYLOGENY

A. Loading Trait Database

In the R code chunk below, do the following:

1. import the raw phosphorus growth data, and
2. standardize the data for each strain by the sum of growth rates.

```
# 1
p.growth <- read.table("./data/p.isolates.raw.growth.txt", sep = "\t", header = TRUE, row.names = 1)

# 2
p.growth.std <- p.growth / (apply(p.growth, 1, sum))
```

B. Trait Manipulations

In the R code chunk below, do the following:

1. calculate the maximum growth rate (μ_{max}) of each isolate across all phosphorus types,
2. create a function that calculates niche breadth (nb), and
3. use this function to calculate nb for each isolate.

```
# 1
umax <- (apply(p.growth, 1, max))

# 2
levins <- function(p_xi) {
  p = 0
  for (i in p_xi) {
    p = p + i^2
  }
  nb = 1 / (length(p_xi) * p)
  return(nb)
}

# 3
nb <- as.matrix(levins(p.growth.std))
rownames(nb) <- row.names(p.growth)
colnames(nb) <- c("NB")
```

C. Visualizing Traits on Trees

In the R code chunk below, do the following:

1. pick your favorite substitution model and make a Neighbor Joining tree,
2. define your outgroup and root the tree, and
3. remove the outgroup branch.

```
# 1
nj.tree <- bionj(seq.dist.F84)

# 2
outgroup <- match("Methanosarcina", nj.tree$tip.label)
nj.rooted <- root(nj.rooted, outgroup, resolve.root = TRUE)
```

```
# 3
nj.rooted <- drop.tip(nj.rooted, "Methanosarcina")
```


In the R code chunk below, do the following:

1. define a color palette (use something other than “YlOrRd”),
2. map the phosphorus traits onto your phylogeny,
3. map the *nb* trait on to your phylogeny, and
4. customize the plots as desired (use `help(table.phylo4d)` to learn about the options).

```
# 1
mypalette <- colorRampPalette(brewer.pal(9, "RdBu"))
```


```
# 2
par(mar = c(1, 1, 1, 1) + 0.1)
x <- phylo4d(nj.rooted, p.growth.std)
```

```
table.phylo4d(x, treetype = "phylo", symbol = "colors", show.node = TRUE, cex.label = 0.5, scale = F, u
```



```
# 3
par(mar = c(1, 5, 1, 5) + 0.1)
x.nb <- phylo4d(nj.rooted, nb)
```

```
table.phylo4d(x.nb, treetype = "phylo", symbol = "colors", show.node = TRUE, cex.label = 0.5, scale = F
```



Question 6:

- Make a hypothesis that would support a generalist-specialist trade-off.
- What kind of patterns would you expect to see from growth rate and niche breadth values that would support this hypothesis?

Answer 6a: Species that can grow well on many resources have a low maximum growth rate.

Answer 6b: Species with a high niche breadth have a low u_{max} , while species that have a low niche breadth have a high u_{max} .

6) HYPOTHESIS TESTING

A) Phylogenetic Signal: Pagel's Lambda


In the R code chunk below, do the following:

- create two rescaled phylogenetic trees using lambda values of 0.5 and 0,
- plot your original tree and the two scaled trees, and
- label and customize the trees as desired.

```
# 1
nj.lambda.5 <- rescale(nj.rooted, "lambda", 0.5)
nj.lambda.0 <- rescale(nj.rooted, "lambda", 0)

# 2
layout(matrix(c(1, 2, 3), 1, 3), width = c(1, 1, 1))
par(mar = c(1, 0, 5, 0.5) + 0.1)
```

```
plot(nj.rooted, main = "lambda = 1", cex = 0.7, adj = 0.5)
plot(nj.lambda.5, main = "lambda = 0.5", cex = 0.7, adj = 0.5)
plot(nj.lambda.0, main = "lambda = 0", cex = 0.7, adj = 0.5)
```



In the R code chunk below, do the following:

1. use the `fitContinuous()` function to compare your original tree to the transformed trees.

```
fitContinuous(nj.rooted, nb, model = "lambda")
```

```
## GEIGER-fitted comparative model of continuous data
```

```
## fitted 'lambda' model parameters:
```

```
## lambda = 0.061979
```

```
## sigsq = 0.140020
```

```
## z0 = 0.664039
```

```
##
```

```
## model summary:
```

```
## log-likelihood = 21.456738
```

```
## AIC = -36.913476
```

```
## AICc = -36.227761
```

```
## free parameters = 3
```

```
##
```

```
## Convergence diagnostics:
```

```
## optimization iterations = 100
```

```
## failed iterations = 57
```

```
## frequency of best fit = NA
```

```
##
```

```
## object summary:
```

```
## 'lik' -- likelihood function
## 'bnd' -- bounds for likelihood search
## 'res' -- optimization iteration summary
## 'opt' -- maximum likelihood parameter estimates

fitContinuous(nj.lambda.0, nb, model = "lambda")

## GEIGER-fitted comparative model of continuous data
## fitted 'lambda' model parameters:
## lambda = 0.000000
## sigsq = 0.139249
## z0 = 0.656203
##
## model summary:
## log-likelihood = 21.399126
## AIC = -36.798252
## AICc = -36.112537
## free parameters = 3
##
## Convergence diagnostics:
## optimization iterations = 100
## failed iterations = 0
## frequency of best fit = 0.93
##
## object summary:
## 'lik' -- likelihood function
## 'bnd' -- bounds for likelihood search
## 'res' -- optimization iteration summary
## 'opt' -- maximum likelihood parameter estimates
```

Question 7: There are two important outputs from the `fitContinuous()` function that can help you interpret the phylogenetic signal in trait data sets. a. Compare the lambda values of the untransformed tree to the transformed (lambda = 0). b. Compare the Akaike information criterion (AIC) scores of the two models. Which model would you choose based off of AIC score (remember the criteria that the difference in AIC values has to be at least 2)? c. Does this result suggest that there's phylogenetic signal?

Answer 7a: Untransformed lambda = 0.062, transformed lambda = 0 **Answer 7b:** The difference in AIC is about 0.12, so the models are about equivalent. **Answer 7c:** Nope.

B) Phylogenetic Signal: Blomberg's K

In the R code chunk below, do the following:

1. correct tree branch-lengths to fix any zeros,
2. calculate Blomberg's K for each phosphorus resource using the `phylosignal()` function,
3. use the Benjamini-Hochberg method to correct for false discovery rate, and
4. calculate Blomberg's K for niche breadth using the `phylosignal()` function.

```
# 1
nj.rooted$edge.length <- nj.rooted$edge.length + 10^-7

# 2
p.phylosignal <- matrix(NA, 6, 18)
colnames(p.phylosignal) <- colnames(p.growth.std)
rownames(p.phylosignal) <- c("k", "PIC.vars.obs", "PIC, var.mean", "PIC.var.P", "PIC.var.z", "PIC.P.BH")
```

```

for (i in 1:18) {      # This calculates Bloomberg's k for each resource
  x <- as.matrix(p.growth.std[, i, drop = FALSE]) # store column "x" as matrix
  out <- phylosignal(x, nj.rooted) # compare column "x" to rooted tree
  p.phylosignal[1:5, i] <- round(t(out), 3) # Round
}

# 3
p.phylosignal[6, ] <- round(p.adjust(p.phylosignal[4, ], method = "BH"), 3)

# 4
signal.nb <- phylosignal(nb, nj.rooted)
signal.nb

```

```

##           K PIC.variance.obs PIC.variance.rnd.mean PIC.variance.P
## 1 4.187392e-06          49966.79          49922.28          0.546
## PIC.variance.Z
## 1      0.002136296

```

Question 8: Using the K-values and associated p-values (i.e., “PIC.var.P”) from the phylosignal output, answer the following questions:

- Is there significant phylogenetic signal for niche breadth or standardized growth on any of the phosphorus resources?
- If there is significant phylogenetic signal, are the results suggestive of clustering or overdispersion?

Answer 8a: There is no significant phylogenetic signal for niche breadth. There is phylogenetic signal for cAMP and DNA ($p < 0.05$) **Answer 8b:** $k=0$, so results suggest overdispersion.

C. Calculate Dispersion of a Trait

In the R code chunk below, do the following:

- turn the continuous growth data into categorical data,
- add a column to the data with the isolate name,
- combine the tree and trait data using the `comparative.data()` function in `caper`, and
- use `phylo.d()` to calculate D on at least three phosphorus traits.

```

# 1
p.growth.pa <- as.data.frame((p.growth > 0.01) * 1)

# 2
p.growth.pa$name <- rownames(p.growth.pa)

# 3
p.traits <- comparative.data(nj.rooted, p.growth.pa, "name")

# 4
phylo.d(p.traits, binvar = DNA)

##
## Calculation of D statistic for the phylogenetic structure of a binary variable
##
## Data : p.growth.pa
## Binary variable : DNA
## Counts of states: 0 = 20

```



```

##              1 = 19
## Phylogeny : nj.rooted
## Number of permutations : 1000
##
## Estimated D : 0.6001446
## Probability of E(D) resulting from no (random) phylogenetic structure : 0.031
## Probability of E(D) resulting from Brownian phylogenetic structure : 0.006
phylo.d(p.traits, binvar = cAMP)

##
## Calculation of D statistic for the phylogenetic structure of a binary variable
##
## Data : p.growth.pa
## Binary variable : cAMP
## Counts of states: 0 = 10
##              1 = 29
## Phylogeny : nj.rooted
## Number of permutations : 1000
##
## Estimated D : 0.1023278
## Probability of E(D) resulting from no (random) phylogenetic structure : 0.001
## Probability of E(D) resulting from Brownian phylogenetic structure : 0.37
phylo.d(p.traits, binvar = ATP)

##
## Calculation of D statistic for the phylogenetic structure of a binary variable
##
## Data : p.growth.pa
## Binary variable : ATP
## Counts of states: 0 = 1
##              1 = 38
## Phylogeny : nj.rooted
## Number of permutations : 1000
##
## Estimated D : 5.931215
## Probability of E(D) resulting from no (random) phylogenetic structure : 0.882
## Probability of E(D) resulting from Brownian phylogenetic structure : 0.065

```

Question 9: Using the estimates for D and the probabilities of each phylogenetic model, answer the following questions:

- Choose three phosphorus growth traits and test whether they are significantly clustered or overdispersed?
- How do these results compare the results from the Blomberg's K analysis?
- Discuss what factors might give rise to differences between the metrics.

Answer 9a: DNA and cAMP are significantly overdispersed, and ATP is marginally overdispersed.

Answer 9b: These are consistent with Blomberg's K analysis. **Answer 9c:** When we calculate D for each metric we aren't correcting for multiple comparisons, so we might expect to see more significantly clumped or overdispersed traits using D than K .

7) PHYLOGENETIC REGRESSION

In the R code chunk below, do the following:

1. Load and clean the mammal phylogeny and trait dataset,
2. Fit a linear model to the trait dataset, examining the relationship between mass and BMR,
2. Fit a phylogenetic regression to the trait dataset, taking into account the mammal supertree

```
# 1
mammal.Tree <- read.tree("./data/mammal_best_super_tree_fritz2009.tre")
mammal.data <- read.table("./data/mammal_BMR.txt", sep = "\t", header = T)


mammal.data <- mammal.data[, c("Species", "BMR_.ml02.hour.", "Body_mass_for_BMR_.gr.")]
mammal.species <- array(mammal.data$Species)

pruned.mammal.tree <- drop.tip(mammal.Tree, mammal.Tree$tip.label[-na.omit(match(mammal.species, mammal

pruned.mammal.data <- mammal.data[mammal.data$Species %in% pruned.mammal.tree$tip.label, ]
rownames(pruned.mammal.data) <- pruned.mammal.data$Species

# 2
fit <- lm(log10(BMR_.ml02.hour.) ~ log10(Body_mass_for_BMR_.gr.), data = pruned.mammal.data)

plot(log10(pruned.mammal.data$Body_mass_for_BMR_.gr.), log10(pruned.mammal.data$BMR_.ml02.hour.), las =
abline(a = fit$coefficients[1], b = fit$coefficients[2])
b1 <- round(fit$coefficients[2], 3)
eqn <- bquote(italic(z) == .(b1))
text(0.5, 4.5, eqn, pos = 4)
```




```
# 3
```

```
fit.phy <- phylolm(log10(BMR_.ml02.hour.) ~ log10(Body_mass_for_BMR_.gr.), data = pruned.mammal.data, p
AIC(fit)
```

```
## [1] -365.5042
```

```
plot(log10(pruned.mammal.data$Body_mass_for_BMR_.gr.), log10(pruned.mammal.data$BMR_.ml02.hour.), las =
abline(a = fit.phy$coefficients[1], b = fit.phy$coefficients[2])
b1.phy <- round(fit.phy$coefficients[2],3)
eqn <- bquote(italic(z) == .(b1.phy))
text(0.5, 4.5, eqn, pos = 4)
```



- Why do we need to correct for shared evolutionary history?
- How does a phylogenetic regression differ from a standard linear regression?
- Interpret the slope and fit of each model. Did accounting for shared evolutionary history improve or worsen the fit?
- Try to come up with a scenario where the relationship between two variables would completely disappear when the underlying phylogeny is accounted for.

Answer 10a: Species on our tree aren't independent - they share evolutionary history. So we need to account for that in our model. **Answer 10b:** The variance of the error term in a phylogenetic regression is described by a covariance matrix that takes into account branch lengths. In a standard linear regression the variance of the error term is a constant (sigma squared).

Answer 10c: The slope of the model that accounts for phylogeny has a higher slope than the simple regression model, meaning the relationship between mass and BMR has a higher slope than we'd think if we didn't account for shared ancestry. The AIC value for model that accounts for phylogeny is much lower than the simple regression model, so it's a better fit. **Answer 10d:** Say you wanted to know how shouting affected plant growth rate, so you had a shouting treatment with 10 tree species and a no-shouting treatment with 10 grass species. At the end of the experiment you found that the shouting treatment had a higher growth rate! But of course that's just because they're all trees, and trees grow faster than grasses. So once you account for shared ancestry, there wouldn't be any effect of shouting (probably...). That's my best shot!

7) SYNTHESIS


Work with members of your Team Project to obtain reference sequences for 10 or more taxa in your study. Sequences for plants, animals, and microbes can be found in a number of public repositories, but perhaps the most commonly visited site is the National Center for Biotechnology Information (NCBI) <https://www.ncbi.nlm.nih.gov/>:

[//www.ncbi.nlm.nih.gov/](http://www.ncbi.nlm.nih.gov/). In almost all cases, researchers must deposit their sequences in places like NCBI before a paper is published. Those sequences are checked by NCBI employees for aspects of quality and given an **accession number**. For example, here an accession number for a fungal isolate that our lab has worked with: JQ797657. You can use the NCBI program nucleotide **BLAST** to find out more about information associated with the isolate, in addition to getting its DNA sequence: <https://blast.ncbi.nlm.nih.gov/>. Alternatively, you can use the `read.GenBank()` function in the `ape` package to connect to NCBI and directly get the sequence. This is pretty cool. Give it a try.

But before your team proceeds, you need to give some thought to which gene you want to focus on. For microorganisms like the bacteria we worked with above, many people use the ribosomal gene (i.e., 16S rRNA). This has many desirable features, including it is relatively long, highly conserved, and identifies taxa with reasonable resolution. In eukaryotes, ribosomal genes (i.e., 18S) are good for distinguishing coarse taxonomic resolution (i.e. class level), but it is not so good at resolving genera or species. Therefore, you may need to find another gene to work with, which might include protein-coding gene like cytochrome oxidase (COI) which is on mitochondria and is commonly used in molecular systematics. In plants, the ribulose-bisphosphate carboxylase gene (*rbcL*), which on the chloroplast, is commonly used. Also, non-protein-encoding sequences like those found in **Internal Transcribed Spacer (ITS)** regions between the small and large subunits of the ribosomal RNA are good for molecular phylogenies. With your team members, do some research and identify a good candidate gene.

After you identify an appropriate gene, download sequences and create a properly formatted fasta file. Next, align the sequences and confirm that you have a good alignment. Choose a substitution model and make a tree of your choice. Based on the decisions above and the output, does your tree jibe with what is known about the evolutionary history of your organisms? If not, why? Is there anything you could do differently that would improve your tree, especially with regard to future analyses done by your team?

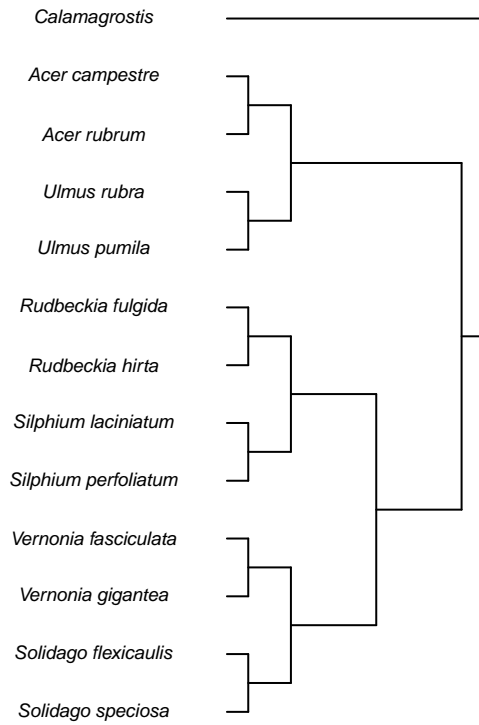
```
# Visualize
read.aln <- read.alignment(file = "./data/plant_seq.afa", format = "fasta")
p.DNABin <- as.DNABin(read.aln)
window <- p.DNABin[, 100:500]
image.DNABin(window, cex.lab = 0.5)
```



```
seq.dist.F84 <-dist.dna(p.DNAbin, model = "F84", pairwise.deletion = FALSE)
F84.tree <-bionj(seq.dist.F84)
F84.outgroup <-match("Calamagrostis", F84.tree$tip.label)
F84.rooted <-root(F84.tree, F84.outgroup, resolve.root = TRUE)

layout(matrix(c(1, 2), 1, 2), width = c(1, 1))
par(mar = c(1, 0, 2, 1))
plot.phylo(F84.rooted, type = "phylogram", direction = "left", show.node.label = TRUE, use.edge.length = TRUE)
```

F84



Yep, the tree jibes with what is known! *Rudbeckia*, *Silphium*, *Vernonia*, and *Solidago* are all Asteraceae, so it makes sense that they form a clade. We chose two species from each genus, and the species grouped together by their genera. Hooray!